

How the redaction service works, and six things we need from you

DATE

Monday 2026-09-28 · 18:00–19:00 ET · Week 6 of 15

GOAL

Walk you through the design across three workstreams, and close six open items — GPU, NDA, midterm, Drive, chat, SOW

A

Who does what

Mark · 4 min

B

How the service works

Zhexuan & Anmol · 18 min

C

Deciding what to redact

Yudi · 20 min

D

Evaluation dataset

Karina · 4 min

E

Six things we need

Mark · 14 min

A · Who does what MARK · 4 MIN

One person per area, matching the five responsibilities Aarvin set out. Confirmed at our team meeting on 2026-09-25.

AREA	OWNER	WHAT IT COVERS
1 · Evaluation & Test Data	Karina Yang	The labelled email benchmark, evaluation metrics, regression testing
2 · Privacy Research	Yudi Guo	What counts as private, the category and tier definitions, regulatory grounding
3 · Model Work	Anmol Sahu	Prompting, running the local model, producing confidence scores

4 • Service Engineering	Zhexuan Ye	The local service itself, the HTTP contract, latency, failure handling
5 • Integration & Coordination	Mark Chu	Documentation, progress updates, the final demo

Mark is your single point of contact, as requested.

B • How the service works ZHEXUAN & ANMOL • 18 MIN

Zhexuan and Anmol's workflow and architecture proposal, in full. It reworks the earlier hybrid routing diagrams against the direction set on Sep 21.

EMBEDDED DOCUMENT • REDACTION-ARCHITECTURE-V2-PROPOSAL • ZHEXUAN YE • 2026-09-27

CAPSTONE • VIRTUAL GOLD INC • PROPOSAL — NOT YET MERGED

Local redaction service — workflow & architecture (v2 proposal)

2026-09-27 • Zhexuan Ye • reworks the diagrams in `hybrid-routing-diagrams-en.html` against Aarvin George's Sep 23 email

In plain terms: before any text reaches an outside AI assistant, it comes through this service first.



- **Whole-request check first.** A separate check looks at the request as a whole before any finer judgment matters: if it falls into a category that is not yet approved for release at all — a confidential thread, health or financial information, an unclear source — the entire request stays local, full stop, and nothing below that check ever sees it.
- **Piece-by-piece judgment, once approved.** A local model reads the cleared request and checks it against a global, cross-session store of known-safe and known-sensitive entities, then decides, span by span, how confident it is that something is private: anything it's confident about gets redacted automatically, and anything it's unsure about is not guessed at — it's queued for a human reviewer instead. That review happens asynchronously: the caller gets an immediate "still deciding" response rather than waiting on a person, and checks back separately once it's resolved.

- **Placeholders only, ever.** Either way, the caller only ever gets text back with the sensitive parts replaced by placeholders; the mapping back to the original words never leaves this machine.
- **Restore is separate and deliberate.** If the caller's own system later needs the original wording — say, after getting a reply from an outside AI assistant — it can ask this service to restore it using a reference id, as a separate step, not something that happens on its own.
- **No outside calls from here.** This service never calls an outside AI assistant itself; that decision belongs entirely to the calling application.

`hybrid-routing-diagrams-en.html` is untouched.

Figure 1 (v2) · Request workflow






Same lane style as the original hybrid-routing diagram, now with the merged-in policy gate and asynchronous review. Steps 1–7 always run: the gateway forwards the request, the classifier consults the local context store and asks the local model to detect spans and judge context, then compiles a sensitivity report for the new **Policy gate**. The policy gate makes one conservative, deterministic release decision for the whole request: if it is a confidential, health, financial or credential category, or the source/provenance is unclear, it returns **LOCAL_ONLY** straight back to the Gateway (**7L**, red dashed) — nothing below it ever sees that request, and no per-request review can override this until the approved policy itself changes. Only an approved request continues (8) to the confidence gate, which branches once on per-span confidence (no more local-vs-cloud split): 8A auto-redacts through Redact / restore, 8B flags the span in the Review queue — which immediately acknowledges **PENDING_REVIEW** (10) without waiting for a human, so the call still returns promptly. A reviewer resolves it out-of-band, at their own pace, and writes the outcome into the same record Redact / restore holds. Either path is logged with no raw text (9), and the result (or the pending acknowledgement) bubbles back up — through the policy gate, then the classifier and gateway lanes (10) — to reach the caller (11). Three separate, id-keyed calls happen later, independent of this main flow and of each other: **restore** (R1–R3, which can itself come back **REVIEW_BLOCKED** if its own leak scan fails), **checking review status** (V1–V3, the only way a caller learns a pending review's outcome and gets the content), and **deleting the record** (D1–D3) — all three reach Redact / restore

directly, bypassing the classifier and both gates. Band E at the bottom runs offline on test data only. Corner icons: sparkle = model, cylinder = data store, </> = hand-written code.

Figure 2 (v2) • Component architecture (numbered)



New: a whole-request policy gate, merged in from a teammate's architecture v0.6 and interim HTTP contract. Their design makes one conservative release decision for the entire request (confidential category, health/financial/credential data, or unknown provenance), separate from the per-span redaction decision our confidence gate already made. So the classifier's report (7) now goes to this new **Policy gate** first: if the whole

request is blocked, it deterministically returns **LOCAL_ONLY** straight to the Gateway (**L**, red dashed), skipping the confidence gate, review queue and redact/restore entirely — nothing below this line ever sees that request, and this can't be resolved per-request, only by changing the approved policy. Only an approved request continues (8) to the confidence gate, which auto-redacts (8A) or flags for asynchronous human review (8B, **PENDING_REVIEW**) exactly as before. Both gates log their decision to the audit trail (7L, 9). When 8B fires, Review queue immediately acks **PENDING_REVIEW** without waiting — a human reviewer resolves it out-of-band, at their own pace, and writes the outcome straight into the record Redact / restore holds (**V0**). Local model stays next to the sensitivity classifier (5, 6); the dashed **R/V/D** connector bundles three later, independent, id-keyed calls that all reach Redact / restore directly, bypassing the classifier and both gates: **restore** (can itself come back **REVIEW_BLOCKED**), **poll for a pending review's outcome** (the only way a caller gets content after **PENDING_REVIEW**), and **delete the record**. Redact / restore's tagline now separates its two kinds of state: the scan/review **record always exists**, but the restorable **mapping is opt-in** — setting ``mapping.retain`` is the only thing that decides whether restore can substitute real values back in. Cloud API and egress proxy remain removed, per the client's framing that the calling application decides on its own, later, whether to call an external cloud assistant. This is a proposal for discussion — [hybrid-routing-diagrams-en.html](#) has not been edited. Figure 1 now shows the same policy-gate and asynchronous-review steps (7, 7L, 8, V0).

Open question 1 — pending client confirmation: policy-gate granularity. The policy gate's release decision is whole-request, not per-span. As merged in from the teammate's design, one confidential line, or one line of health/financial/credential data, or an unclear source anywhere in the thread is enough to keep the *entire* email local (**LOCAL_ONLY**), even if the rest of the message is ordinary business text — and because this is a deterministic policy decision, it can't be resolved by a per-request human review the way a low-confidence span can.

- This is likely to trigger often in practice — long threads and quoted history make it easy for one sensitive line to bounce back a whole email that would otherwise be fine to release.
- The teammate's doc treats this as a deliberate, conservative interim default (see architecture v0.6, "Apply the interim policy"), not a final design — the source/type/task policy table and category boundaries are explicitly still open for the client privacy/legal representative and system owner to approve.
- **Recommendation:** raise this with the client as a concrete open item — whether finer-grained rules (e.g. per-paragraph or per-quoted-message release, rather than per whole thread) are worth the added complexity, or whether the coarse whole-request block is acceptable for the first release.

Open question 2 — pending team decision: how the async review queue is checked, and how long it waits. 8B now returns `PENDING_REVIEW` immediately and a human resolves it later (V0); this fixes the earlier problem of a synchronous call implying a blocking human wait, but it opens three mechanical questions the diagram currently answers with a single default.

- **Poll vs. webhook.** This proposal has the caller poll `GET /v1/redactions/{id}` (V1–V3) until it resolves. A webhook/callback the caller registers would avoid repeated polling but adds delivery/retry/auth surface of its own — worth deciding before the API contract is final.
- **Review TTL.** How long can a request sit as `PENDING_REVIEW` before it expires? The diagram assumes it shares the same bounded-lifetime record as everything else in Redact / restore, but that cap was originally sized for restore/scan state, not for however long a human reviewer actually takes — these may need separate limits.
- **Cancel via delete.** D1–D3 assumes deleting a still-pending review simply cancels it. If the client instead wants pending reviews to be un-cancelable once queued (e.g. for audit completeness), that changes D's behavior.

Local redaction service · workflow & architecture proposal v2 · Virtual Gold capstone · 2026-09-27

YOUR FEEDBACK

Where the contract or the flow does not match your system

OWNER

Zhexuan

FOLLOW-UP

date

C · Deciding what to redact YUDI · 20 MIN

Yudi's sensitive data classification proposal, in full, below. **Read it against Karina's redact-always design** — that one masks every span any detector finds and restores locally, which sets the ceiling on safety and the floor on usefulness. **It is the baseline: everything below has to earn its place against it by measurement, not by argument.** Karina's document is in `docs/research/karina/2026-09-24/`.

CMU CAPSTONE × VIRTUAL GOLD · PROPOSAL

Sensitive Data Classification Proposal

SCOPE

How the local service determines which content to redact, and how the approach will be validated

PREPARED BY

Yudi Guo · Privacy research & evaluation

DATE

27 Sep 2026

STATUS


Draft proposal; not yet validated on data

Overview

Each detected item is assessed on two dimensions: its **category**, which determines the severity of a potential leak, and its **context**, which determines the confidence that it is private. A **rule-based decision layer** combines the two. Where confidence is insufficient, the item is redacted by default and replaced with a reversible placeholder, so the assistant’s task is unaffected.

REQUIREMENT	APPROACH
Local execution on an open-source model	All detectors and the model run on project hardware. The service has no outbound network access.
Contextual judgment by the model	The model assesses each candidate span within its surrounding text, for example “my cell” versus “our support line”.
Confidence score output	The model scores each span; the score is calibrated against labelled emails before use.

System Architecture



The model produces assessments (step 03); decisions are made by code (step 04) against a client-approved table. All components other than the cloud assistant run within the local service.

Technical Feasibility IMPLEMENTATION

Every stage is built on existing open-source tools. The only substantial custom component is the merging of detector outputs.

COMPONENT	IMPLEMENTATION	EFFORT
Rule-based detection	Regular expressions with checksum validation for phone numbers, email addresses, card and ID numbers. Presidio includes a set of such recognisers.	Low
Presidio	The <code>presidio-analyzer</code> package with a spaCy language model [2]	Low

Local model	Served through Ollama with Rescriber's open-source detection prompt [1]; OpenAI Privacy Filter as an alternative [3]	Low to medium
Merging outputs (union)	Custom code that aligns results by character position; overlapping spans are counted once	Medium, approx. 50–100 lines

INTEGRATION CONSIDERATIONS

ISSUE	RESOLUTION
The model returns text rather than character positions	Each returned string is located in the source text. Rescriber's prompt requires verbatim substrings for this reason [1].
Detectors use different category names	All outputs are mapped to a single category list. Rescriber's <code>presidio_server.py</code> provides a reference mapping [1].
Latency accumulates across detectors	Detectors run in parallel. The model is the bottleneck: on an Apple M2 machine with 24 GB of memory, Rescriber reports a mean of 0.73 s for Presidio and 3.40 s for Llama3-8B [1].

Sensitivity Tiers FOR CLIENT REVIEW

Each category is assigned to a tier. The tier determines the confidence required before redaction: the more severe the potential harm, the lower the threshold.

TIER	CATEGORIES	REDACTION RULE
HIGH	Government IDs, passwords and keys, bank and card numbers, health information	Redact at confidence ≥ 0.2
HIGH	Business-confidential information: unreleased financial results, deal and M&A terms, project code names, legal matters	Redact at confidence ≥ 0.2 ; see below

MEDIUM

Names of private individuals, personal phone numbers, personal email addresses, home addresses

Redact at confidence ≥ 0.5

LOW

City, dates, organisation names, age, gender

Retain when isolated; redact when several items together could identify an individual

PUBLIC

Public figures, published support lines, company names in a public context

Retain only when the model is confident the item is public

BUSINESS-CONFIDENTIAL INFORMATION

This category cannot be listed exhaustively and has no fixed format, so rules and Presidio do not detect it. Coverage relies on three layers:

LAYER	MECHANISM	COVERS
1. Term list	The client supplies project code names, deal names and key counterparties; any match is redacted	Confidential items the client already knows
2. Explicit markers	Phrases such as “confidential”, “board only”, “do not forward” or “privileged” place the surrounding passage in the High tier	Content the sender has marked as confidential
3. Default placeholder	Proper nouns and amounts the model cannot classify are replaced with reversible placeholders	Items missed by the first two layers

KNOWN LIMITATION

NIST SP 800-122 covers personal data only, so tiering for business information will follow the client’s internal data classification. This category depends on the model’s contextual judgment and is the principal recall risk; its leak rate will be reported separately. **PROPOSED**

✓ BASIS FOR THE TIERS AND THRESHOLD SELECTION**BASIS FOR THE TIERS**

NIST SP 800-122 classifies personal data by low, moderate and high confidentiality impact, taking into account the sensitivity of each field and its context of use, and notes that combinations of fields may be more sensitive than individual fields [5]. The Low-tier

combination rule reflects Sweeney's finding that ZIP code, gender and date of birth together uniquely identify an estimated 87% of the US population [6]. The category list is adapted from Rescriber's taxonomy [1].

THRESHOLD SELECTION

The thresholds above are initial values. Final values will be determined empirically:

1. The client sets an acceptable leak rate for each tier. This is a business decision.
2. A range of thresholds is evaluated on labelled emails, recording the leak rate and the over-redaction rate at each.
3. For each tier, the highest threshold that meets its leak-rate target is selected, which minimises unnecessary redaction.

THRESHOLD	LEAK RATE	OVER-REDACTION RATE
0.2	1%	40%
0.3	2%	25%
0.5	5%	12%
0.7	12%	5%

Illustrative values for the Medium tier. With a 2% leak-rate target, the selected threshold would be 0.3.

VALIDATION AND SIGN-OFF

Given the dataset size (100 emails), thresholds will be tuned and evaluated on separate subsets, or by cross-validation, to avoid optimistic results. Tier assignments are a business decision and require client confirmation.

Worked Example

The example includes the case raised in the client's email: a personal mobile number and a public support number in the same message.

INCOMING · FROM DANA WHITFIELD TO TOM REYES

Hi Tom,

Before Thursday’s board call: my oncologist moved my appointment, so I’ll dial in from home. Reach me on my cell, 412-555-0182.

On the vendor outage, Acme’s support line is 1-800-555-0199.

— Dana

EXECUTIVE’S REQUEST TO THE ASSISTANT

Draft a reply to Dana and summarise the vendor issue.

SPAN	CATEGORY · TIER	CONTEXTUAL EVIDENCE	CONFIDENCE (PRIVATE)	DECISION
Dana, Dana Whitfield	Name · MED	Private individual; sender	0.90	[PERSON_1]
Tom	Name · MED	Private individual; recipient	0.85	[PERSON_2]
oncologist	Health · HIGH	“my ... appointment”	0.95	[HEALTH_1]
412-555-0182	Phone · MED	“my cell”	0.95	[PHONE_1]
1-800-555-0199	Phone · MED	Toll-free prefix; “Acme’s support line”	0.15	Retain
Acme, Thursday	Organisation, date · LOW	Isolated; insufficient to identify an individual	n/a	Retain

Confidence values are illustrative.

TEXT SENT TO THE CLOUD ASSISTANT

Hi [PERSON_2],

Before Thursday's board call: my [HEALTH_1] moved my appointment, so I'll dial in from home. Reach me on my cell, [PHONE_1].

On the vendor outage, Acme's support line is 1-800-555-0199.

— [PERSON_1]

OUTCOME

The assistant can still draft the reply to [PERSON_1] and summarise the Acme outage. The cloud service does not receive the diagnosis, the personal number or either name. Placeholders are restored locally before the draft is returned to the executive.

Confidence Scoring

The model scores each span. Because language models tend to be overconfident [7], the raw score is calibrated against labelled emails before thresholds are applied.

✓ CALIBRATION PROCEDURE AND ALTERNATIVE METHOD

VERSION 1: MODEL SCORE WITH CALIBRATION

The model assigns a raw score to each span. Language models tend to be overconfident when stating their own confidence [7], so the raw score is calibrated against labelled emails before use:

1. Run the system on labelled emails, recording each span's raw score and its true label.
2. Group spans by raw score and compute the proportion in each group that is actually private.
3. At run time, convert each raw score to the corresponding observed proportion.

RAW SCORE	SPANS	ACTUALLY PRIVATE	CALIBRATED SCORE
0.9–1.0	200	140	0.70
0.7–0.9	100	50	0.50
0.5–0.7	50	15	0.30
0.0–0.5	150	15	0.10

Illustrative values. A raw score of 0.95 would be reported as 0.70.

Calibration corrects a raw score; it is not a separate source of scores. It ensures that a threshold of 0.5 corresponds to an observed 50% likelihood.

VERSION 2: AGREEMENT-BASED SCORING

If calibration shows that the model’s scores discriminate poorly, for example with most spans scored near 0.9, the raw score will instead be derived from agreement: the share of detectors that flag a span, and the share of 3–5 repeated model runs that label it private. Agreement across repeated answers has been shown to be a more reliable uncertainty signal than stated confidence [7]. Both methods will be compared on the same data.

Related Work and Contributions

No individual component is novel. Each has a published source or an industry precedent.

CONTRIBUTIONS OF THIS PROJECT

1. Integration of these components for executive email, without per-item human review. Rescriber, the closest prior system, requires the user to confirm each item.

2. A labelled email benchmark and an ablation study quantifying the contribution of each component.

✓ SOURCES FOR EACH COMPONENT

COMPONENT	SOURCE	REPORTED FINDING OR CAPABILITY
Local detection, placeholders sent to the cloud, restoration of the response	Rescriber, CHI 2025 [1]	On 240 test samples: Llama3-8B precision 0.94, recall 0.63; GPT-4o recall 0.88; Presidio 0.69 / 0.68. MIT-licensed code.
Rule-based and NER detection; allow-lists	Microsoft Presidio [2]	Widely used open-source toolkit; <code>allow_list</code> excludes known-safe values
Local PII detection model	OpenAI Privacy Filter [3]	Open-weight, Apache 2.0, deployable on-premises

Per-category thresholds	Google Cloud Sensitive Data Protection [4]	Five likelihood levels; minimum likelihood configurable per data type
Sensitivity tiers	NIST SP 800-122 [5]	Low / moderate / high impact levels; context of use can change the level
Automated decisions from scores and thresholds	MINIM, ICML 2026 [8]; OCELOT, 2026 [9]	Scored, rule-based decisions in place of human review

Improving Recall PRINCIPAL RISK

Rescriber’s 8B model missed approximately one third of personal data (recall 0.63) [1]. Its users reviewed every message; here, a missed item reaches the cloud directly. A baseline will first be measured on email, and the measures below added one at a time.

MEASURES AND SUPPORTING EVIDENCE

MEASURE	RATIONALE	EVIDENCE
Union of all detectors	The union’s recall is at least that of its best member; the cost is additional over-redaction.	HOLDS BY DEFINITION
Entity propagation	Once a name is detected, every occurrence in the text is marked. Rescriber reports missed variants of the same name.	FAILURE MODE IN [1]
Repeated model runs	Spans flagged in any run become candidates.	HOLDS BY DEFINITION
Email-specific prompt examples	Rescriber’s prompt examples are chat messages, not email threads.	PROPOSED
One category per pass	A shorter category list per pass may reduce omissions.	UNTESTED

Evaluation Plan

100 realistic emails, as requested, with each private span labelled by category and by whether it is private or public in context. Public labelled datasets will supplement this set [10]. The primary metric is the **leak rate**: the proportion of private spans sent to the cloud, reported per tier.

FULL METRIC SET AND ABLATION STUDY

METRICS

METRIC	DEFINITION
Leak rate (primary)	Proportion of private spans sent to the cloud, reported per tier
Over-redaction rate	Proportion of public spans redacted unnecessarily
Calibration	Agreement between calibrated scores and observed proportions
Latency	Processing time per email on the confirmed hardware

ABLATION STUDY

Starting from a Rescriber-style baseline, components are added one at a time: detector union → entity propagation → repeated runs → tier thresholds. Thresholds are tuned on one subset and results are reported on a held-out subset.

Failure Handling

The service fails safe: if any component fails, every candidate span is redacted.

FAILURE-HANDLING RULES

CONDITION	BEHAVIOUR
Model unavailable or output unparseable	Fall back to rule-based detection, redact all candidates, and flag the response as degraded

Response contains a placeholder that was not issued	Flag the placeholder rather than infer a value. Rescriber observed GPT-4o generating placeholders such as “[Your Full Name]” [1].
Logging	Record category and score only; original text is never logged
Model and list updates	Installed at deployment; the running service does not access the internet

Next Steps WITH THE CLIENT

The following items will be finalised with the client once the prototype is demonstrated, using measured results rather than assumptions.

ITEM	HOW IT WILL BE DETERMINED	TIMING
Input scope	Plain text only, or also attachments and scanned documents. This defines what the prototype must handle.	Before prototyping
Confidential term list	The client supplies project code names, deal names and key counterparties for the business-confidential term list.	Before prototyping
Tier assignments	The prototype’s redactions on sample emails are reviewed with the client, who adjusts the tier table accordingly.	After the demo
Thresholds	The client sets an acceptable leak rate per tier; thresholds are then selected from experimental results.	After the demo
Latency	Processing time is measured on the confirmed hardware and compared with the client’s expectations.	After the demo
Optional context fields	Over-redaction is compared with and without fields such as internal or external sender, to show whether they are worth providing.	After the demo

✓ SOURCES (VERIFIED 24–27 SEP 2026; FIGURES ARE TAKEN FROM EACH SOURCE’S OWN EXPERIMENTS)

1. Zhou, Xu, Wu, Li. Rescriber: Smaller-LLM-Powered User-Led Data Minimization for LLM-Based Chatbots. CHI 2025. Code: github.com/PEACH-Research-Lab/Rescriber
2. Microsoft Presidio, allow-lists
3. OpenAI Privacy Filter (April 2026)
4. Google Cloud Sensitive Data Protection, match likelihood
5. NIST SP 800-122, Guide to Protecting the Confidentiality of PII (2010), §3.2
6. Sweeney. Simple Demographics Often Identify People Uniquely. CMU Data Privacy Working Paper 3, 2000
7. Xiong et al. Can LLMs Express Their Uncertainty? ICLR 2024
8. Yu et al. MINIM: Privacy-Aware Minimal View for Agents via Trusted Local Sanitization. ICML 2026
9. Xie, Li. OCELOT: Inference-Leakage Budgets for Privacy-Preserving LLM Agents. 2026
10. Agyemang et al. LLM-Redactor. 2026. Benchmark on Hugging Face: [jayluxferro/llm-redactor-leak-benchmark](https://huggingface.co/jayluxferro/llm-redactor-leak-benchmark)

YOUR FEEDBACK

Objections, corrections, what the tiers and thresholds should be

OWNER

Yudi

FOLLOW-UP

date

D • Evaluation dataset KARINA • 4 MIN

A short introduction only. The detail goes in its own document.

The 100 labelled emails Aarvin asked for: where they come from, what gets recorded on each one, and how we check the labels are trustworthy before scoring anything against them.

YOUR FEEDBACK

Notes on the dataset

OWNER

Karina

FOLLOW-UP

date

E • Six things we need from you

MARK • 14 MIN

Ordered by what blocks us soonest.

1 • GPU resources

What hardware will be available, and when?

We confirmed on Sep 27 that the team has **personal laptops only**. That is workable for the first phase — the detection layer runs on CPU, so **we can give you a working endpoint without any GPU at all**. The second phase, where the local model does the contextual judgement, needs a T4-class card or better.

The sandbox offered on Sep 3 has been an open question since our Sep 15 meeting: provider, GPU, direct access, and who pays. **It is the only thing blocking one of the five workstreams entirely.**

2 • NDA process

How is the NDA addendum progressing?

Building a standalone service with an HTTP endpoint, rather than integrating directly, was explicitly the workaround while this is pending. Two things turn on it:

- Whether the design stays as it is, or changes once we can see your architecture
- **Whether the API request can tell us where each piece of text came from** — body, quoted thread, the user's own words. We only need that coarse level, not anything about your internal structures, and it would measurably reduce our error rate. It may not be blocked by the NDA at all, which is worth checking.

3 • Midterm presentation — Monday 2026-10-05

We would like the midterm to be a working session with you, not a one-way presentation.

What we want to close on the day:

TO CLOSE	WHY THEN
The project definition, as a final version	The direction changed on Sep 21 and has been settling since. The midterm is the point to fix it rather than keep revising
Resources, confirmed	GPU and sandbox, so the second phase is not still blocked halfway through the semester

The Scope of Work, signed

See item 6 — we revise it from what is agreed today, and ask you to sign it on the day

The timeline from there to the end

We bring it; we would like your agreement on the milestones rather than just informing you of them

We will also show what is running by then, with measurements of what redaction costs in answer quality.

4 • Google Drive

Confirm the shared Drive and who has access.

From the Sep 21 meeting: we use Virtual Gold's Drive, with access for Aarvin, Alex, Inderpal and Urte. We can set this up straight after this call if you confirm the list.

5 • Google Chat

Are we moving to Google Chat?

The Slack workspace showed a 30-day trial expiry on our accounts, and a single workspace covering two capstones as separate threads needs a paid feature. Moving to Google Chat was raised as the alternative. We are happy either way — we just need one place, decided.

6 • Scope of Work

What we propose: we revise the Scope of Work from whatever is agreed in today's meeting, send it to you in advance, and ask you to finalise and sign it at the midterm on 10/05.

That way it is signed against the work as it actually stands, not as it stood in week 3.

ANSWERS AND DECISIONS

GPU · NDA · midterm · Drive · chat · SOW

OWNER

Mark

FOLLOW-UP

date